

# classification

November 6, 2025

```
[1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

```
[2]: pop = pd.read_csv("pop_genres_dataset.csv")  
pop.head()
```

```
[2]:      artists track_name  popularity duration_ms explicit \
0   my little airport           47     183968       0
1   my little airport           44     172343       0
2   my little airport           50     191308       0
3   my little airport           48     148564       0
4    Eason Chan                54     228973       0

      danceability   energy   key  loudness   mode speechiness acousticness \
0        0.576    0.505    11   -13.070     0      0.0392      0.753
1        0.577    0.228    11   -14.684     1      0.0306      0.834
2        0.641    0.479     4   -15.395     1      0.0443      0.732
3        0.613    0.210     5   -15.309     1      0.0309      0.769
4        0.646    0.370     0   -10.980     1      0.0351      0.825

  instrumentalness  liveness  valence    tempo time_signature track_genre
0      0.199000     0.151    0.570  127.025             4    cantopop
1      0.000805     0.384    0.161  119.035             4    cantopop
2      0.059800     0.115    0.301  121.974             4    cantopop
3      0.000477     0.139    0.268  132.035             4    cantopop
4      0.000005     0.090    0.465  129.914             4    cantopop
```

## 0.1 Predicting pop song explicitness

```
[3]: from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score,  
recall_score, ConfusionMatrixDisplay, roc_curve, roc_auc_score  
from sklearn.model_selection import train_test_split, cross_val_score, KFold
```

```
[4]: # Only selecting numeric variables  
X = pop[["popularity", "duration_ms", "danceability", "energy",  
         "key", "loudness", "mode", "speechiness", "acousticness",
```

```

    "instrumentalness", "liveness", "valence", "tempo", "time_signature"]]
y = pop["explicit"]

# Splitting into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42)

```

[5]: df = pd.concat([X\_train, y\_train], axis = 1)  
df.corr()

	popularity	duration_ms	danceability	energy	key	\
popularity	1.000000	0.119770	0.133854	0.019746	0.011389	
duration_ms	0.119770	1.000000	-0.028426	-0.092777	0.006021	
danceability	0.133854	-0.028426	1.000000	0.054455	0.031049	
energy	0.019746	-0.092777	0.054455	1.000000	0.018415	
key	0.011389	0.006021	0.031049	0.018415	1.000000	
loudness	0.105342	-0.111904	0.050184	0.720586	0.024051	
mode	-0.071275	-0.020384	-0.120570	-0.091753	-0.092927	
speechiness	0.068684	-0.112708	0.131527	0.226773	0.046018	
acousticness	0.044592	0.135789	-0.116178	-0.708240	-0.012636	
instrumentalness	-0.045519	-0.056019	-0.060166	0.054256	-0.013831	
liveness	-0.054945	-0.058741	-0.129208	0.140642	-0.025546	
valence	-0.014605	-0.070367	0.334176	0.463023	0.027881	
tempo	-0.055654	-0.048555	-0.210335	0.179279	0.012591	
time_signature	0.040957	-0.013938	0.131912	0.193845	0.019084	
explicit	0.014291	-0.092259	0.129864	0.034781	0.018659	

	loudness	mode	speechiness	acousticness	\
popularity	0.105342	-0.071275	0.068684	0.044592	
duration_ms	-0.111904	-0.020384	-0.112708	0.135789	
danceability	0.050184	-0.120570	0.131527	-0.116178	
energy	0.720586	-0.091753	0.226773	-0.708240	
key	0.024051	-0.092927	0.046018	-0.012636	
loudness	1.000000	-0.054929	0.164364	-0.487139	
mode	-0.054929	1.000000	-0.073293	0.077547	
speechiness	0.164364	-0.073293	1.000000	-0.157578	
acousticness	-0.487139	0.077547	-0.157578	1.000000	
instrumentalness	-0.058715	0.022493	-0.019572	-0.077869	
liveness	0.043576	0.002541	0.060020	-0.099658	
valence	0.211999	-0.066680	0.138043	-0.315047	
tempo	0.109209	0.024638	0.122292	-0.146461	
time_signature	0.146760	-0.013669	0.037423	-0.210545	
explicit	0.079670	-0.016466	0.151343	-0.105286	

	instrumentalness	liveness	valence	tempo	\
popularity	-0.045519	-0.054945	-0.014605	-0.055654	
duration_ms	-0.056019	-0.058741	-0.070367	-0.048555	

danceability	-0.060166	-0.129208	0.334176	-0.210335
energy	0.054256	0.140642	0.463023	0.179279
key	-0.013831	-0.025546	0.027881	0.012591
loudness	-0.058715	0.043576	0.211999	0.109209
mode	0.022493	0.002541	-0.066680	0.024638
speechiness	-0.019572	0.060020	0.138043	0.122292
acousticness	-0.077869	-0.099658	-0.315047	-0.146461
instrumentalness	1.000000	0.015621	0.001061	0.010143
liveness	0.015621	1.000000	0.043661	0.023649
valence	0.001061	0.043661	1.000000	0.078323
tempo	0.010143	0.023649	0.078323	1.000000
time_signature	-0.002642	0.006178	0.088197	-0.009007
explicit	-0.008930	0.022528	-0.012827	0.012941

	time_signature	explicit
popularity	0.040957	0.014291
duration_ms	-0.013938	-0.092259
danceability	0.131912	0.129864
energy	0.193845	0.034781
key	0.019084	0.018659
loudness	0.146760	0.079670
mode	-0.013669	-0.016466
speechiness	0.037423	0.151343
acousticness	-0.210545	-0.105286
instrumentalness	-0.002642	-0.008930
liveness	0.006178	0.022528
valence	0.088197	-0.012827
tempo	-0.009007	0.012941
time_signature	1.000000	0.031345
explicit	0.031345	1.000000

In the training data, explicitness seems to have low correlations with the predictor variables. As such, Random Forests should be used to investigate non-linear relationships and combinations that can help predict whether or not a song is explicit.

### 0.1.1 Cross validation

```
[6]: # Getting the best number of estimators and depth
estimators = np.arange(100, 300, 50)
depths = np.arange(1, 21)
best_score = 0
best_depth = 0
best_estimator = 0

fig,ax = plt.subplots(1)
ax.set(xlabel = "Depth", ylabel = "Cross validations core")

for e in estimators:
```

```

for d in depths:
    rf = RandomForestClassifier(n_estimators= 100, max_depth = d)

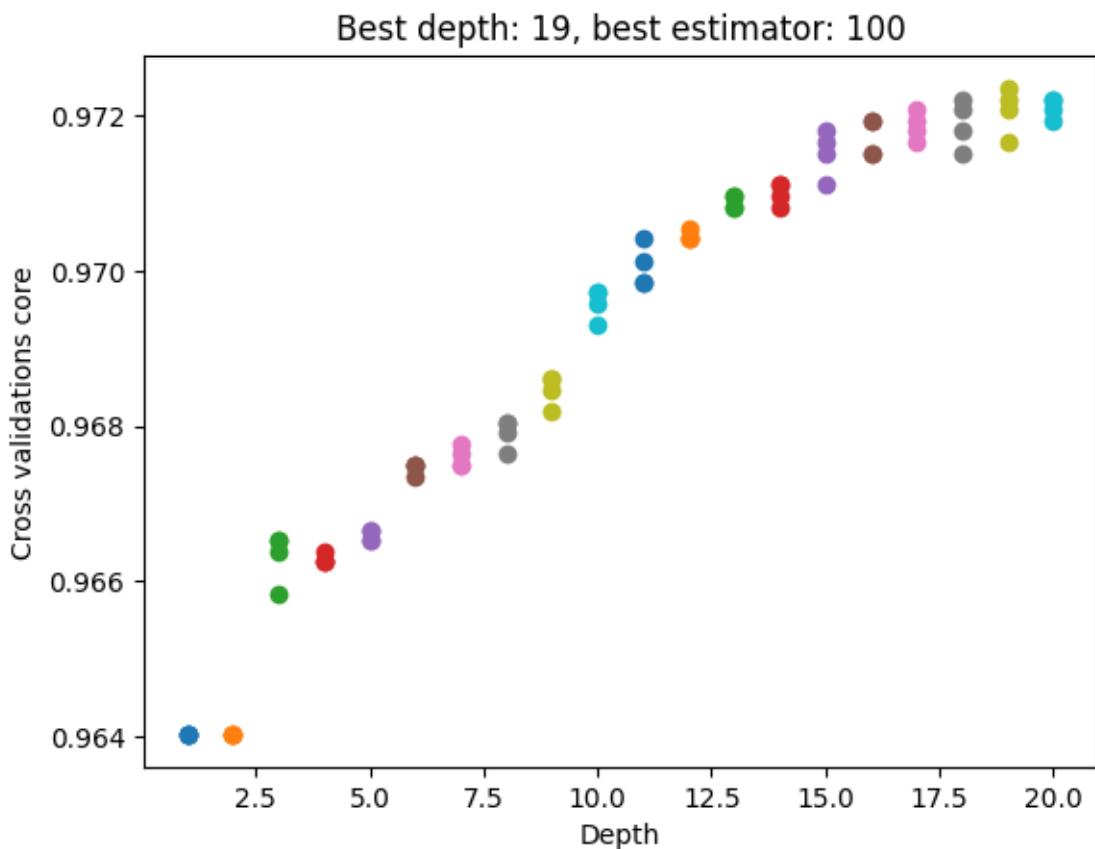
    # Cross validate with 5 folds
    score = cross_val_score(rf, X_train, y_train, cv = 5, u
    ↪scoring="accuracy").mean()

    if score >= best_score:
        best_score = score
        best_depth = d
        best_estimator = e

    ax.scatter(d, score)

ax.set(title = f"Best depth: {best_depth}, best estimator: {best_estimator}")
plt.show()

```



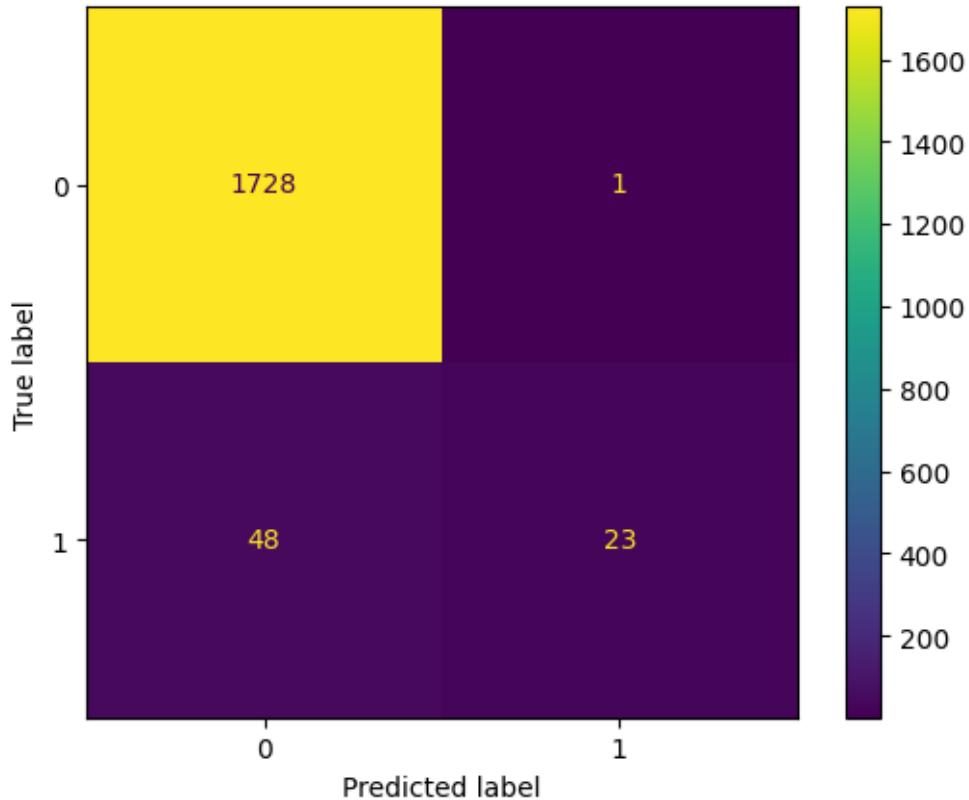
## 0.1.2 Random Forests

```
[7]: rf = RandomForestClassifier(n_estimators= best_estimator, max_depth =  
    ↪best_depth)  
rf.fit(X_train, y_train)  
y_pred = rf.predict(X_test)
```

## 0.1.3 Confusion matrix

```
[8]: ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
```

```
[8]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x10d532150>
```



## 0.1.4 Performance metrics

```
[9]: print("Prediction accuracy:", round(accuracy_score(y_test, y_pred), 5))  
print("Prediction error:", round(1 - accuracy_score(y_test, y_pred), 5))  
  
# To obtain values in matrix  
cm = confusion_matrix(y_test, y_pred)  
tn, fp, fn, tp = cm.ravel()
```

```

print("True positive rate:", round(tp / (tp + fn), 5))
print("True negative rate:", round(tn / (tn + fp), 5))

precision = tp / (tp + fp)
recall = tp / (tp + fn)

print("F1 score:", round((2 * precision * recall) / (precision + recall), 5))

```

Prediction accuracy: 0.97278  
 Prediction error: 0.02722  
 True positive rate: 0.32394  
 True negative rate: 0.99942  
 F1 score: 0.48421

Overall, the model predicts the data well with a prediction accuracy of 97.33%. The model is excellent at identifying non-explicit songs 99.94% of the time but is only able to identify explicit songs 33.8% of the time. This may be due to the fact that there are significantly more non-explicit songs than explicit songs in the data. The F1 score supports the fact that the model is unable to predict explicit songs very well.

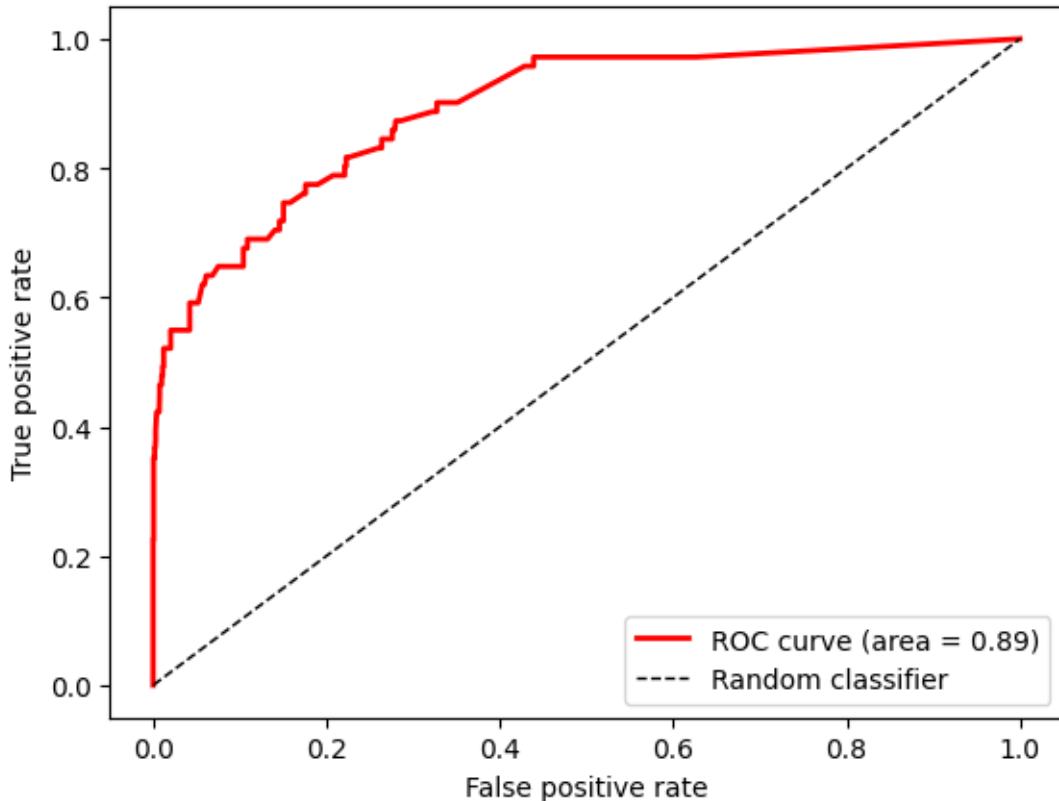
### 0.1.5 ROC curve

```

[10]: y_pred_proba = rf.predict_proba(X_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

plt.figure()
plt.plot(fpr, tpr, color = "red", lw = 2,
          label=f"ROC curve (area = {roc_auc:.2f})")
plt.plot([0,1], [0,1], color = "black", lw = 1, linestyle = "--",
          label = "Random classifier")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.legend(loc="lower right")
plt.show()

```



After observing the ROC curve with the AUC being 0.89, the problem may lie in the model's threshold instead of its ability to discern between explicit vs non explicit. Adjusting the threshold may lead to better classification of explicit songs.

```
[ ]: kf = KFold(n_splits=5, shuffle=True, random_state=42)
auc_per_fold = cross_val_score(rf, X, y, cv=kf, scoring="roc_auc")
accuracy_per_fold = cross_val_score(rf, X, y, cv=kf, scoring="accuracy")
fold = pd.DataFrame({
    "Fold": np.arange(1, 6),
    "AUC": auc_per_fold,
    "Accuracy": accuracy_per_fold
})
print(fold)
```

	Fold	AUC	Accuracy
0	1	0.904484	0.972778
1	2	0.886606	0.972222
2	3	0.902769	0.975000
3	4	0.898703	0.975556
4	5	0.893236	0.972778